

FSFFL NEXT

Routed Forecast - Final Resolution

Outcome A: Routing Supported | Research freeze only | 18 Sep 2026

The final research result supports a small routing exception rather than using B2a for every player. The exception is narrow: keep B2a almost everywhere, but use the already-developed R2 treatment for certain established or veteran quarterbacks. No model was retrained, refit, promoted, merged, or deployed.

What changed from the prior Outcome C

The prior stop was a reproducibility problem, not a football conclusion. The missing historical row-level artifacts were recovered from the FSFFL NEXT Library and matched the durable repository evidence. That closed Stage 0. This continuation then used the recovered 4,287-row pre-holdout coordinate without rerunning Forecast.

The frozen routing rule

Player group	Y2 treatment	Y3 treatment
QB - developmental (0-3 years experience)	B2a	B2a
QB - established (4-8 years)	R2	R2
QB - veteran (9+ years)	R2	B2a
RB / WR / TE - all career states	B2a	B2a

Career state is determined only from NFL experience known at the forecast cutoff: 0-3 years, 4-8 years, or 9+ years. The study did not use future decline, current-player identities, or named-player outcomes to decide which route a player receives.

Why quarterback career state matters

A simple QB-wide rule looks attractive at first because the richer treatment improves the position average. But that average hides a strong split. R2 is consistently better for established quarterbacks and for veteran quarterbacks at Y2, while it is consistently worse for developmental quarterbacks.

QB cell	Early gain	Mid gain	Validation gain	Result
Developmental Y2	-7.93	-6.07	-10.64	Keep B2a
Developmental Y3	-1.23	-10.49	-15.26	Keep B2a
Established Y2	+6.84	+0.48	+12.09	Use R2
Established Y3	+5.57	+1.62	+4.65	Use R2
Veteran Y2	+22.50	+27.24	+15.41	Use R2
Veteran Y3	+17.82	+30.38	+4.70	Too sparse -> B2a

The numbers are MAE gains versus B2a, so positive is better. Veteran QB Y3 is directionally favorable, but the mid and validation samples fall below the pre-frozen 20-row support floor. It therefore defaults to B2a.

How much does the frozen route improve the full historical sample?

Metric	Universal B2a	Frozen route	Change
MAE	41.7557	41.0709	-0.6847
Total absolute error	179,006.6	176,071.2	-2,935.4
Mean signed error	-15.4315	-13.2574	2.1741 closer to zero
Y2 MAE	41.0893	40.0077	-1.0816
Y3 MAE	42.5842	42.3929	-0.1913

Across all 4,287 legal pre-holdout observations, MAE improves by about 1.64%. The paired historical uncertainty range for the MAE gain is +0.36 to +1.02 points. More importantly, the improvement is not coming from one fold: every early, mid, and validation point estimate improves at both Y2 and Y3.

Why the extra career-state split is worth the complexity

The best full-chronology position-only challenger routes every QB to R2 at Y2 and leaves everything else on B2a. That gets overall MAE to 41.4020. The frozen position x career-state route reaches 41.0709, another 0.3311-point improvement. Its paired historical range versus the position-only version is +0.05 to +0.62 points.

The practical reason is clearer than the aggregate number: the position-only rule would worsen developmental-QB Y2 MAE from about 75.50 to 83.86. The final route avoids that damage completely while keeping the gains in established and veteran quarterbacks.

Why R2, not R1 or H2

R2 beats R1 on the selected quarterback cells, including the combined paired evidence. H2 sometimes improves further on its mid/validation rows, but its incremental advantage over R2 is not stable enough to justify the added complexity, and H2 has no leakage-free early-fold evaluation. The simpler R2 treatment is therefore frozen.

Protected cohorts

Cohort	Universal B2a MAE	Frozen route MAE	Interpretation
Top 10% source production (n=479)	60.1136	59.6726	Small improvement
Developmental 0-3 years (n=2,877)	39.9204	39.9204	Exactly unchanged

The important negative: collapse cases get worse

The routing result is not uniformly better. On players who later experience role loss or a severe production collapse, the richer R2 treatment is less conservative and therefore misses high when the collapse occurs. This is a real cost and is preserved in the final result rather than hidden.

Stress cohort	B2a MAE	Frozen route MAE	Change	Share of total error
Role loss (n=586)	34.7706	36.3230	+1.5524	11.38% -> 12.09%
Deep collapse (n=295)	50.6173	53.7254	+3.1081	8.34% -> 9.00%

That stress result does not reverse the management outcome because the governing question is whether a route can be assigned from information known at the forecast cutoff. Earlier FSFFL research found that most of B2a's safety advantage over R2 was concentrated in severe realized collapses and could not establish that those events were reliably knowable in advance. A separate source-time role/security correction was tested and made broad calibration worse.

So the final route does not use future collapse to decide who receives B2a. Instead it reports the cost explicitly: non-role-loss MAE improves from 42.86 to 41.82, non-deep-collapse MAE improves from 41.10 to 40.14, and total absolute error still falls by roughly 2,935 points.

Holdout status

No holdout was used to discover or select this route. After the routing specification was frozen, the repository's holdout checkpoint was rechecked. The one-time final holdout had already been permanently consumed by an earlier frozen candidate study, so there is no untouched final holdout legally available for this routing study. The previously inspected 2021-2022 block is also non-pristine and was not used to rescue the result. No current-player board was run.

Final management conclusion

Outcome A - Routing Supported. Preserve B2a as the default Forecast treatment. Use R2 only for established quarterbacks at Y2/Y3 and veteran quarterbacks at Y2. Keep developmental quarterbacks and every RB/WR/TE state on B2a. This specification is frozen for management review only; it is not implementation or promotion authority.

Governance and reproducibility

Recovered coordinate SHA-256: 8095788e1b48028c50b65d7fba51ac48a41daceeaa6459ae3a219723ce30b1eb.

Reproducible metric-evaluation runtime: 1.681452 seconds.

Forecast model reruns: 0. Forecast model refits: 0. Holdout routing runs: 0. Current-player routing runs: 0.

No changes to main, PR #147, production authority, Intrinsic/Shapley, the Constitution, or global C=0.25.